


```
import pandas as np
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

## UPLOAD DATASET

```
# Import necessary libraries
from google.colab import files
import pandas as pd
```

```
# Upload the CSV file
uploaded = files.upload()
```


 Choose Files Housing (2).csv

- **Housing (2).csv**(text/csv) - 29981 bytes, last modified: 5/15/2025 - 100% done

## DATA EXPLORATION

```
# Read the uploaded CSV file into a pandas DataFrame
# Replace 'your_file.csv' with the actual name of the uploaded file
df = pd.read_csv(next(iter(uploaded)))
```

```
# Display the DataFrame
print(df.head())
```



	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	\
0	13300000	7420	4	2	3	yes	no	no	
1	12250000	8960	4	4	4	yes	no	no	
2	12250000	9960	3	2	2	yes	no	yes	
3	12215000	7500	4	2	2	yes	no	yes	
4	11410000	7420	4	1	2	yes	yes	yes	


  

	hotwaterheating	airconditioning	parking	prefarea	furnishingstatus
0	no	yes	2	yes	furnished
1	no	yes	3	no	furnished
2	no	no	2	yes	semi-furnished
3	no	yes	3	yes	furnished
4	no	yes	2	no	furnished

## DESCRIBE

```
# Get descriptive statistics of the DataFrame
description = df.describe()
```

```
# Display the summary statistics
print(description)
```



	price	area	bedrooms	bathrooms	stories	\
count	5.450000e+02	545.000000	545.000000	545.000000	545.000000	
mean	4.766729e+06	5150.541284	2.965138	1.286239	1.805505	
std	1.870440e+06	2170.141023	0.738064	0.502470	0.867492	
min	1.750000e+06	1650.000000	1.000000	1.000000	1.000000	
25%	3.430000e+06	3600.000000	2.000000	1.000000	1.000000	
50%	4.340000e+06	4600.000000	3.000000	1.000000	2.000000	
75%	5.740000e+06	6360.000000	3.000000	2.000000	2.000000	
max	1.330000e+07	16200.000000	6.000000	4.000000	4.000000	

	parking
count	545.000000
mean	0.693578
std	0.861586
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	3.000000

## DATA CLEANING

### MISSING VALUE

```
# Check for missing values in the DataFrame
missing_values = df.isnull().sum()
```

```
# Display the missing values per column
print(missing_values)

# Optionally, you can visualize the percentage of missing values per column
missing_percentage = (df.isnull().sum() / len(df)) * 100
print(missing_percentage)
```

```
price      0
area       0
bedrooms   0
bathrooms  0
stories    0
mainroad   0
guestroom  0
basement   0
hotwaterheating  0
airconditioning  0
parking     0
prefarea    0
furnishingstatus  0
dtype: int64
price      0.0
area       0.0
bedrooms   0.0
bathrooms  0.0
stories    0.0
mainroad   0.0
guestroom  0.0
basement   0.0
hotwaterheating  0.0
airconditioning  0.0
parking     0.0
prefarea    0.0
furnishingstatus  0.0
dtype: float64
```

## DUPLICATE

```
# Check for duplicate rows
duplicate_rows = df[df.duplicated()]

# Display the number of duplicate rows
print(f"Number of duplicate rows: {duplicate_rows.shape[0]}")

# Optionally, display the duplicate rows themselves
print("\nDuplicate rows:\n")
print(duplicate_rows)
```

```
Number of duplicate rows: 0

Duplicate rows:

Empty DataFrame
Columns: [price, area, bedrooms, bathrooms, stories, mainroad, guestroom, basement, hotwaterheating, airconditioning, parking, prefarea, furnishingstatus]
Index: []
```

## ENCODING

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from google.colab import files

# Handle categorical variables with One-Hot Encoding
df_encoded = pd.get_dummies(df)

# Define features and target
# Automatically detect the target as the last column (or set manually if known)
X = df_encoded.iloc[:, :-1] # All columns except the last
y = df_encoded.iloc[:, -1] # Last column as target

# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Build and train the Linear Regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Predict and evaluate the model
```

```

y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("Model Performance:")
print(f"Mean Squared Error (MSE): {mse:.2f}")
print(f"R2 Score: {r2:.2f}")

```

➞ Model Performance:  
Mean Squared Error (MSE): 0.00  
R<sup>2</sup> Score: 1.00

## MODEL

```

import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error, r2_score

# Assuming y_test and y_pred are already defined from model prediction

```

```

# Evaluation metrics
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

# Print results
print("Model Evaluation Metrics:")
print(f"Mean Squared Error (MSE): {mse:.2f}")
print(f"Root Mean Squared Error (RMSE): {rmse:.2f}")
print(f"R2 Score: {r2:.2f}")


```

```

# Plot Actual vs Predicted
plt.figure(figsize=(8,6))
plt.scatter(y_test, y_pred, alpha=0.6)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], color='red')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.title('Actual vs Predicted')
plt.grid(True)
plt.show()

```

➞ Model Evaluation Metrics:  
Mean Squared Error (MSE): 0.00  
Root Mean Squared Error (RMSE): 0.00  
R<sup>2</sup> Score: 1.00



```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split

```

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from google.colab import files

# Basic data cleaning (drop rows with missing values)
df.dropna(inplace=True)

# One-hot encode categorical variables
df_encoded = pd.get_dummies(df)

# Define features and target
# Replace 'Price' with your actual target column if different
target_column = 'Price'

# : Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train the model
model = LinearRegression()
model.fit(X_train, y_train)

# : Predict
y_pred = model.predict(X_test)

# Evaluate
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

print("Model Evaluation:")
print(f"Mean Squared Error (MSE): {mse:.2f}")
print(f"Root Mean Squared Error (RMSE): {rmse:.2f}")
print(f"R2 Score: {r2:.2f}")

↗ Model Evaluation:
Mean Squared Error (MSE): 0.00
Root Mean Squared Error (RMSE): 0.00
R2 Score: 1.00

# Visualization - Actual vs Predicted
plt.figure(figsize=(8,6))
plt.scatter(y_test, y_pred, alpha=0.6, color='blue')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], color='red')
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title("Actual vs Predicted Prices")
plt.grid(True)
plt.show()
```



## VISUALIZATION

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Optional: improve visual style
sns.set(style="whitegrid")
```

## HISTOGRAM


```
# Print the column names to verify the existence and spelling of 'Price'
print(df.columns)
```

```
# Then proceed with the visualization once you confirm the correct column name
plt.figure(figsize=(10,6))
# Replace 'Price' below with the actual name of the price column from the output of df.columns
# Use the correct column name based on the output of df.columns
sns.histplot(df['price'], kde=True) # Assuming 'price' is the correct column name from the global variable output
plt.title('Distribution of Housing Prices')
plt.xlabel('Price')
plt.ylabel('Frequency')
plt.show()
```

```
Index(['price', 'area', 'bedrooms', 'bathrooms', 'stories', 'mainroad',
      'guestroom', 'basement', 'hotwaterheating', 'airconditioning',
      'parking', 'prefarea', 'furnishingstatus'],
      dtype='object')
```

Distribution of Housing Prices

SCATTER PLOT



```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
plt.figure(figsize=(10,6))
# **IMPORTANT:** Replace 'Area' and 'Price' below with the exact column names from the output of print(df.columns)
# For example, if the output showed ['...','area', 'price', '...'], use sns.scatterplot(x='area', y='price', data=df)
# Assuming from the previous cell's comment that 'price' (lowercase) might be the target column name,
# and assuming 'Area' should also be lowercase based on common naming conventions.
sns.scatterplot(x='area', y='price', data=df) # Changed to lowercase 'area' and 'price' based on common potential issue
plt.title('Price vs Area')
plt.xlabel('Area') # You can keep the label capitalized for readability in the plot
plt.ylabel('Price') # You can keep the label capitalized for readability in the plot
plt.show()
```

